

## Table of Contents

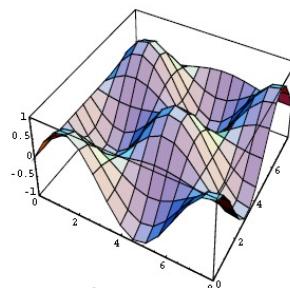
- Motivation & Trends in HPC
- Mathematical Modeling
- Numerical Methods used in HPSC
  - Automatic Differentiation
  - Systems of Differential Equations: ODEs & PDEs
  - Solving Optimization Problems
  - Solving Nonlinear Equations
  - Basic Linear Algebra, Eigenvalues and Eigenvectors
  - Chaotic systems
- HPSC Program Development/Enhancement: from Prototype to Production
- **Visualization, Debugging, Profiling, Performance Analysis, Optimization, Load Balancing**



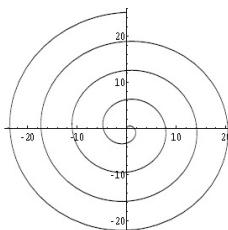
249

## Visualization – Mathematica Graphics

- `Plot3D[Sin[x]*Cos[y], {x,0,8},{y,0,8}]`



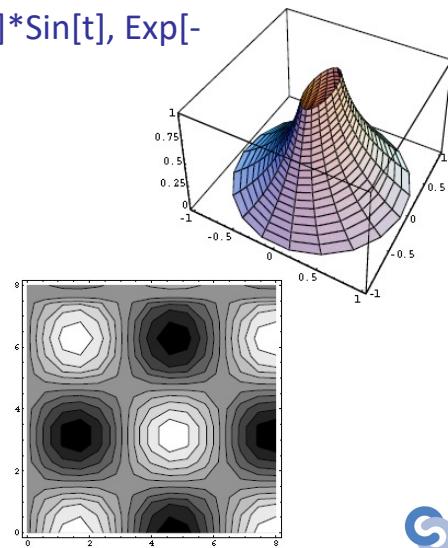
- `ParametricPlot[{t*Sin[t],t*Cos[t]},{t,0,8Pi}]`



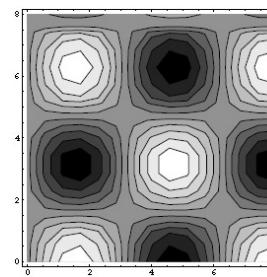
250

## Visualization – Mathematica Graphics

- ParametricPlot3D[{Exp[-2\*u]\*Sin[t], Exp[-u]\*Cos[t], u}, {t, 0, 2Pi}, {u, 0, 1}]



- ContourPlot[Sin[x] Cos[y], {x, 0, 8}, {y, 0, 8}]



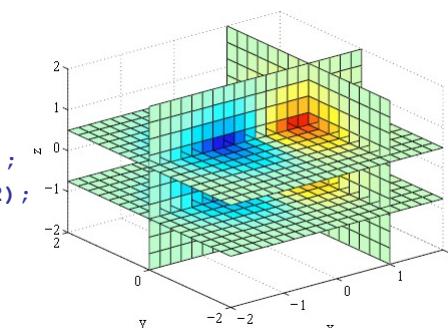
251

## Visualization – Matlab Graphics

- 3D volume visualization of the function  $f(x, y, z) = xe^{-x^2-y^2-z^2}$ , where  $x, y, z \in [-2, 2]$

- Matlab code:

```
r1 = -2:.2:2;
r2 = -2:.25:2;
r3 = -2:.4:2;
[x,y,z] = meshgrid(r1, r2, r3);
v = x.*exp(-x.^2 - y.^2 - z.^2);
colormap(jet);
brighten(0.5);
slice(r1,r2,r3,v, ...
[1.0],[0.0],[-0.75,0.5])
xlabel('x')
ylabel('y')
zlabel('z')
```



252

## Visualization – Matlab Graphics

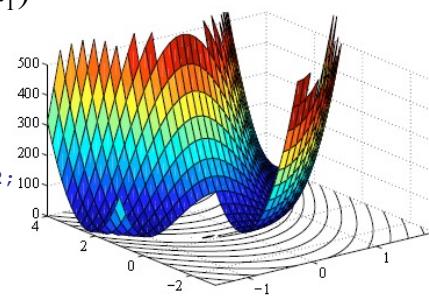
- Surface and contour plot of the Rosenbrock function  $100(x_2 - x_1^2)^2 + (1-x_1)^2$

- Matlab code:

```

x1 = -1.5:0.1:1.9;
x2 = -3:0.2:4;
[xx1 xx2] = meshgrid(x1, x2);
z = 100*(xx2-xx1.^2).^2 + (1-xx1).^2;
nan = NaN;
z0 = z;
% removal of points:
z0(z0 > 600) = nan*z0(z0 > 600);
z0(1:20,1:15) = nan(ones(20,15));
hold on
% surface plot:
surf(x1,x2,z0)
set(gcf, 'DefaultLineLineWidth', 2)
contour(xx1,xx2,z,(0:1.4:50).^3);
hold off
axis([-1.5 1.9 -3 4 0 500])
caxis([0,500])

```



## Foster's Design Methodology

- From *Designing and Building Parallel Programs* by Ian Foster
- Four Steps:
  - **Partitioning**
    - Dividing computation and data
  - **Communication**
    - Sharing data between computations
  - **Agglomeration**
    - Grouping tasks to improve performance
  - **Mapping**
    - Assigning tasks to processors/threads



## Parallel Algorithm Design: PCAM

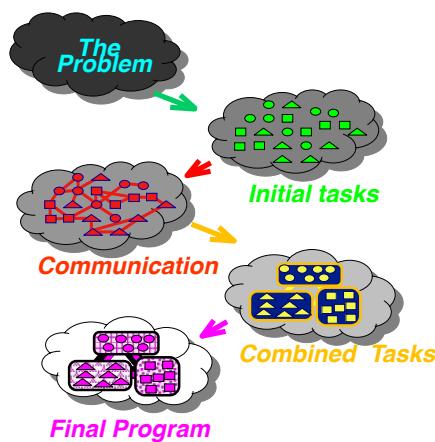
- **Partition:** Decompose problem into fine-grained tasks to maximize potential parallelism
- **Communication:** Determine communication pattern among tasks
- **Agglomeration:** Combine into coarser-grained tasks, if necessary, to reduce communication requirements or other costs
- **Mapping:** Assign tasks to processors, subject to tradeoff between communication cost and concurrency



255

## Designing Threaded Programs

- **Partition**
  - Divide problem into tasks
- **Communicate**
  - Determine amount and pattern of communication
- **Agglomerate**
  - Combine tasks
- **Map**
  - Assign agglomerated tasks to created threads



256

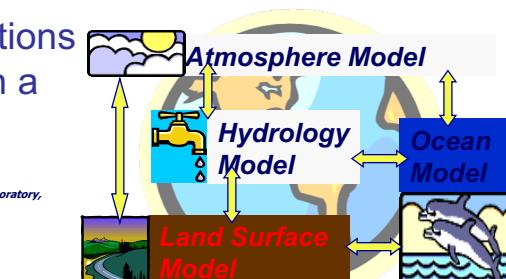
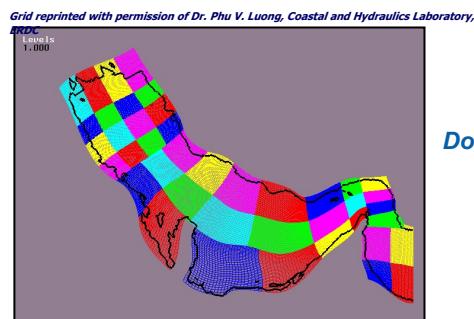
## Parallel Programming Models

- Functional Decomposition
  - Task parallelism
  - Divide the computation, then associate the data
  - Independent tasks of the same problem
- Data Decomposition
  - Same operation performed on different data
  - Divide data into pieces, then associate computation



## Decomposition Methods

- Functional Decomposition
  - Focusing on computations can reveal structure in a problem



- Focus on largest or most frequently accessed data structure
- Data Parallelism
  - Same operation applied to all data



## Example: Computing Pi



- We want to compute  $\pi$
- One method: method of darts\*
- Ratio of area of square to area of inscribed circle proportional to  $\pi$



\*Disclaimer: this is a **TERRIBLE** way to compute  $\pi$ . Don't even think about doing it this way in real life!!!



## Method of Darts

- Imagine dartboard with circle of radius  $R$  inscribed in square
- Area of circle =  $\pi R^2$
- Area of square =  $(2R)^2 = 4R^2$
- Area of circle / Area of square =  $\frac{\pi R^2}{4R^2} = \frac{\pi}{4}$



261

## Method of Darts

- So, ratio of areas proportional to  $\pi$
  - How to find areas?
    - Suppose we threw darts (**completely randomly**) at dartboard
    - Could count number of darts landing in circle and total number of darts landing in square
    - Ratio of these numbers gives approximation to ratio of areas
    - Quality of approximation increases with number of darts
- $-\pi = 4 \times \frac{\# \text{ darts inside circle}}{\# \text{ darts thrown}}$



261

262

## Method of Darts

- Okay, but how in the world do we simulate this experiment on a computer?
  - Decide on length  $R$
  - Generate pairs of random numbers  $(x, y)$  so that  $-R \leq x, y \leq R$
  - If  $(x, y)$  within circle (i.e. if  $(x^2+y^2) \leq R^2$ ), add one to tally for inside circle
  - Lastly, find ratio



262

## Parallelization Strategies

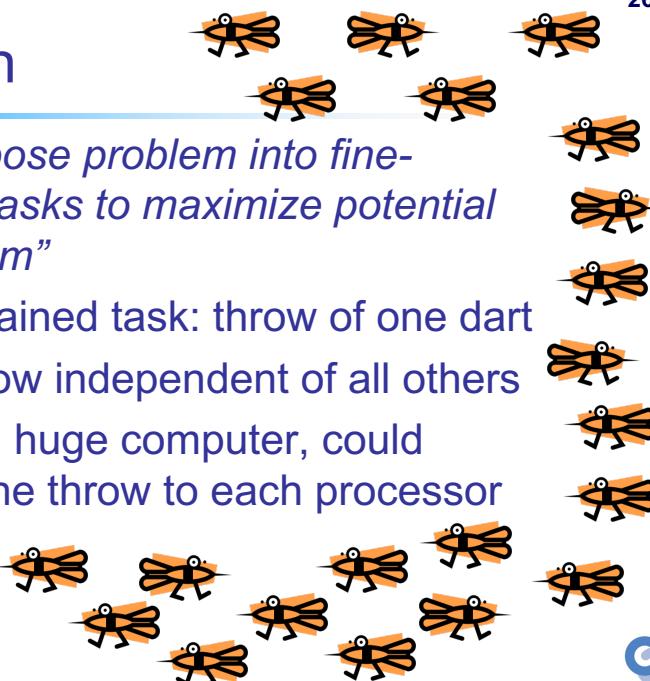
- What tasks independent of each other?
- What tasks must be performed sequentially?
- Using PCAM parallel algorithm design strategy



263

## Partition

- “*Decompose problem into fine-grained tasks to maximize potential parallelism*”
- Finest grained task: throw of one dart
- Each throw independent of all others
- If we had huge computer, could assign one throw to each processor

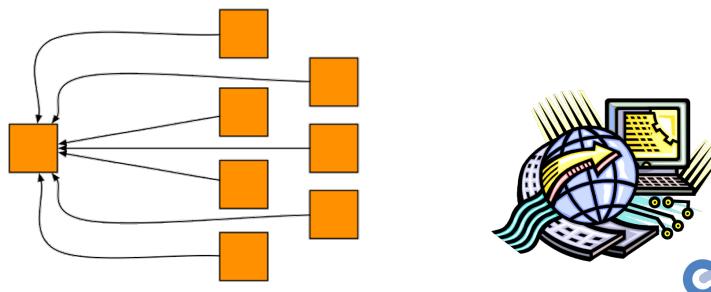


264

## Communication

*“Determine communication pattern among tasks”*

- Each processor throws dart(s) then sends results back to manager process



265

## Agglomeration

*“Combine into coarser-grained tasks, if necessary, to reduce communication requirements or other costs”*

- To get good value of  $\pi$ , must use millions of darts
- We don't have millions of processors available
- Furthermore, communication between manager and millions of worker processors would be very expensive
- Solution: divide up number of dart throws evenly between processors, so each processor does a share of work



266

## Mapping

*“Assign tasks to processors, subject to tradeoff between communication cost and concurrency”*

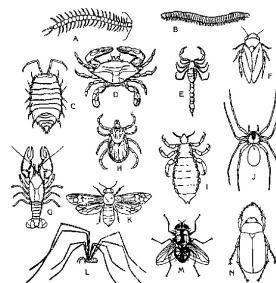
- Assign role of “manager” to processor 0
- Processor 0 will receive tallies from all the other processors, and will compute final value of  $\pi$
- Every processor, including manager, will perform equal share of dart throws



267

## Debugging and Performance Evaluation

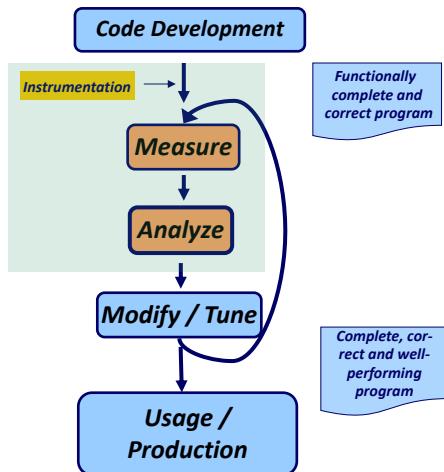
- Common errors in parallel programs
- Debugging tools
- Overview of benchmarking and performance measurements



268

## Concepts and Definitions

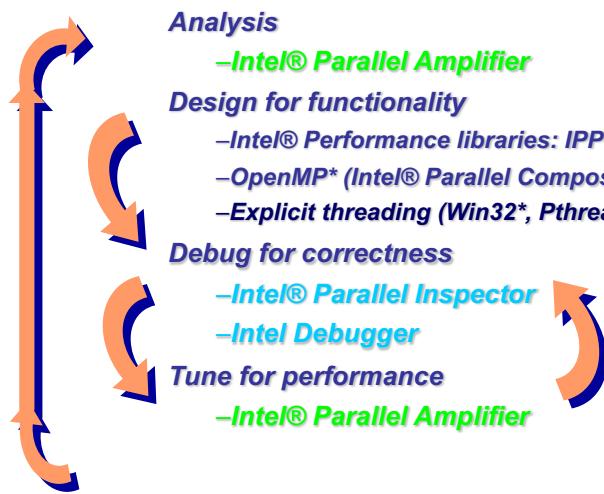
- The typical performance optimization cycle



269



## Development Cycle



270

## Intel® Parallel Studio

- Decide where to add the parallelism
  - Analyze the serial program
  - Prepare it for parallelism
  - Test the preparations
- Add the parallelism
  - Threads, OpenMP, Cilk, TBB, etc.
- Find logic problems
  - Only fails sometimes
  - Place of failure changes
- Find performance problems



271

## Workflow

- Transforming many serial algorithms into parallel form takes five easy high-level steps
- Often existing algorithms are over-constrained by serial language semantics, and the underlying mathematics has a natural parallel expression if you can just find it

(Advisor toolbar)



**Advisor Workflow**

Run [Survey analysis](#) to identify the functions and loops where your program spends most of its time. These are possible places to add parallelism.

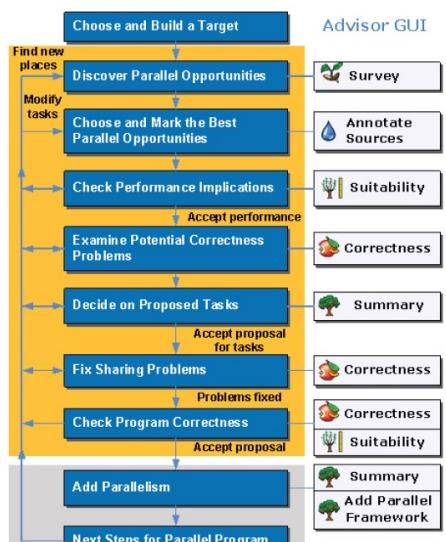
- 1. Survey Target**  
Where should I consider adding parallelism? Locate the loops and functions where your program spends its time, and functions that call them.  
[Start](#) [Explain](#)
- 2. Annotate Sources**  
Add Advisor annotations to [propose](#) parallel tasks and their enclosing parallel sites.  
[Explain](#)
- 3. Check Suitability**  
Analyze the annotated parallel sites and tasks to check their predicted [performance implications](#).  
[Update](#)
- 4. Check Correctness**  
Predict data sharing problems for the annotated tasks and, [fix](#) the reported data sharing problems.  
[Start](#) [Explain](#)
- 5. Add Parallel Framework**  
After you fix problems and re-check your sources, replace Advisor annotations with parallel [framework code](#).  
[Explain](#)

Current Project: 2\_nqueens\_annotated

272

## Advisor Overview

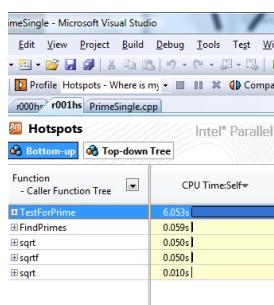
- If you look at these steps in more detail, you find decisions you will need to make
- You do not have to choose the perfect answer the first time, so you can go back and modify your choices



273

## Hotspot Analysis

- Use Parallel Amplification to find hotspots in applications



```

bool TestForPrime(int val)
{
    // let's start checking from 3
    int limit, factor = 3;
    limit = (long)(sqrtf((float)val)+0.5f);
    while( (factor <= limit) && (val % factor) )
        factor++;

    return (factor > limit);
}

void FindPrimes(int start, int end)
{
    // start is always odd
    int range = end - start + 1;
    for( int i = start; i <= end; i+= 2 ){
        if( TestForPrime(i) )
            globalPrimes[gPrimesFound++] = i;
        ShowProgress(i, range);
    }
}
  
```

**Identifies the time consuming regions**

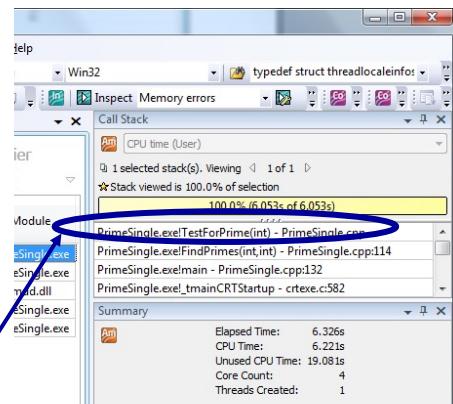


274

## Analysis - Call Stack

- Inspect the code for the leaf node
- Look for a loop to parallelize
  - If none are found, progress up the call stack until you find a suitable loop or function call to parallelize (FindPrimes)

*This is the level in the call tree where we need to thread*

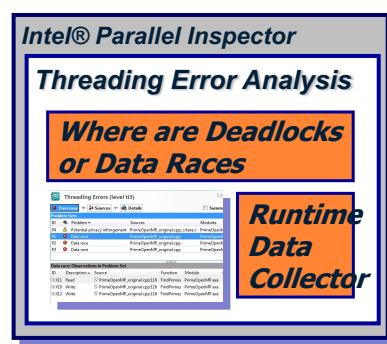


Used to find proper level in the call-tree to thread



## Debugging for Correctness

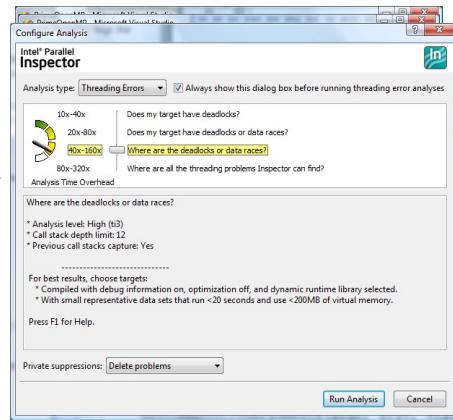
- Intel® Parallel Inspector pinpoints notorious threading bugs like data races and deadlocks



277

## Intel® Parallel Inspector

- Select info regarding both data races & deadlocks
- View the Overview for Threading Errors
- Select a threading error and inspect the code



277

278

## Motivation

- Developing threaded applications can be a complex task
- New class of problems are caused by the interaction between concurrent threads
  - Data races or storage conflicts
    - More than one thread accesses memory without synchronization
  - Deadlocks
    - Thread waits for an event that will never happen



278

279

## Intel® Parallel Inspector

- Debugging tool for threaded software
  - Plug-in to Microsoft\* Visual Studio\*
- Finds threading bugs in OpenMP\*, Intel® Threading Building Blocks, and Win32\* threaded software
- Locates bugs quickly that can take days to find using traditional methods and tools
  - Isolates problems, not the symptoms
  - Bug does **not** have to occur to find it!



279

280

## Parallel Inspector: Analysis

- Dynamic as software runs
  - Data (workload) -driven execution
- Includes monitoring of:
  - Thread and Sync APIs used
  - Thread execution order
    - Scheduler impacts results
  - Memory accesses between threads

**Code path must be executed to be analyzed**



280

## Parallel Inspector: Before You Start

- Instrumentation: background
  - Adds calls to library to record information
    - Thread and Sync APIs
    - Memory accesses
  - Increases execution **time** and **size**
- Use **small** data sets (workloads)
  - Execution time and space is **expanded**
  - Multiple runs over different paths yield best results

**Workload selection is important!**



## Workload Guidelines

- Execute problem code once per thread to be identified
- Use smallest possible working data set
  - Minimize data set size
    - Smaller image sizes
  - Minimize loop iterations or time steps
    - Simulate minutes rather than days
  - Minimize update rates
    - Lower frames per second

**Finds threading errors faster!**



## Binary Instrumentation

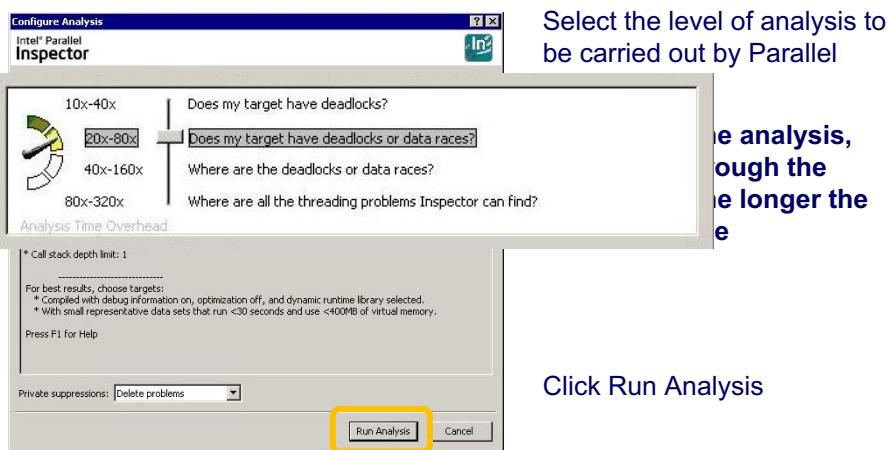
- Build with supported compiler
- Running the application
  - Must be run from within Parallel Inspector
  - Application is instrumented when executed
  - External DLLs are instrumented as used



283

## Starting Parallel Inspector

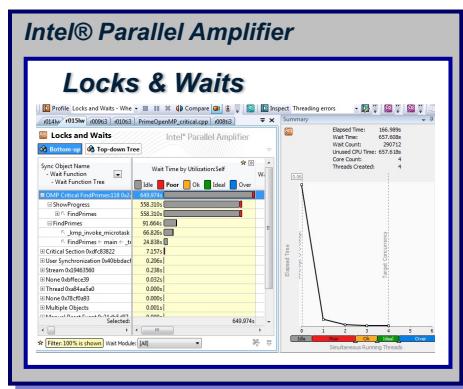
- The Configure Analysis window pops up



284

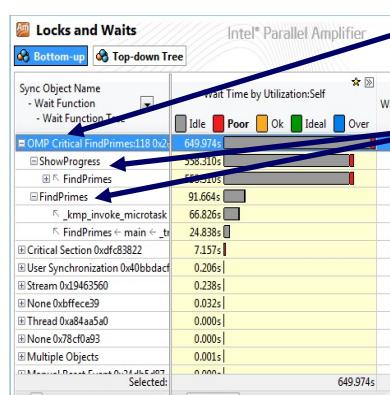
## Tuning for Performance

Parallel Amplifier (Locks & Waits) pinpoints performance bottlenecks in threaded applications



285

## Parallel Amplifier - Locks & Waits



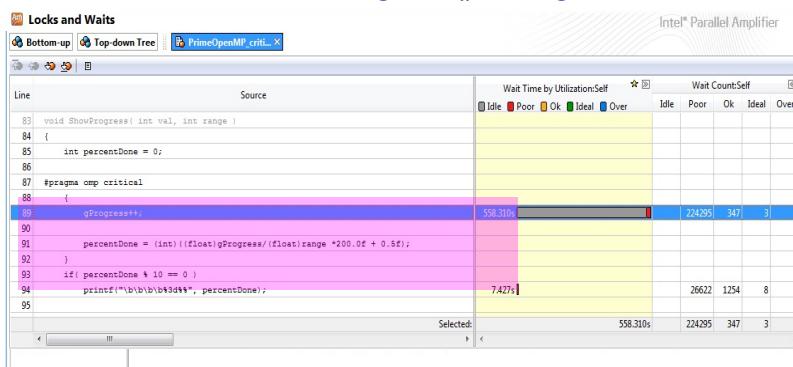
- Graph shows significant portion of time in idle condition as result of critical section
- FindPrimes() & ShowProgress() are both excessively impacted by the idle time occurring in the critical section



286

## Parallel Amplifier - Locks & Waits

- ShowProgress() consumes 558/657 (85%) of the time idling in a critical section
- Double Click ShowProgress() in largest critical section to

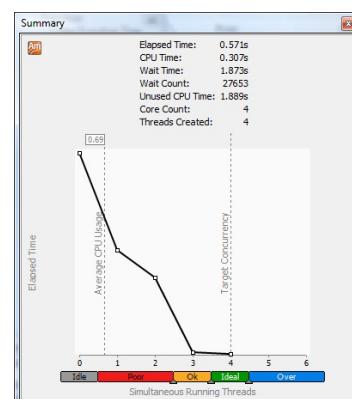


287



## Parallel Amplifier Summary

- Elapsed Time shows .571 sec
- Wait Time/ Core Count =  $1.87/4 = .47$  sec
- Waiting 82% of elapsed time in critical section
- Most of the time 1 core and occasionally 2 are occupied

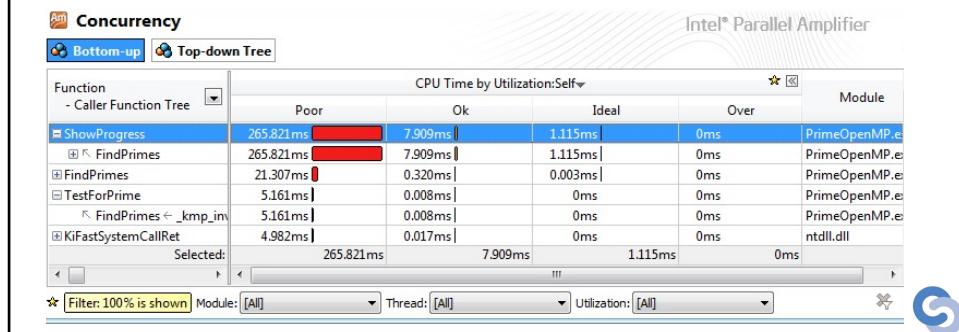


288



## Parallel Amplifier - Concurrency

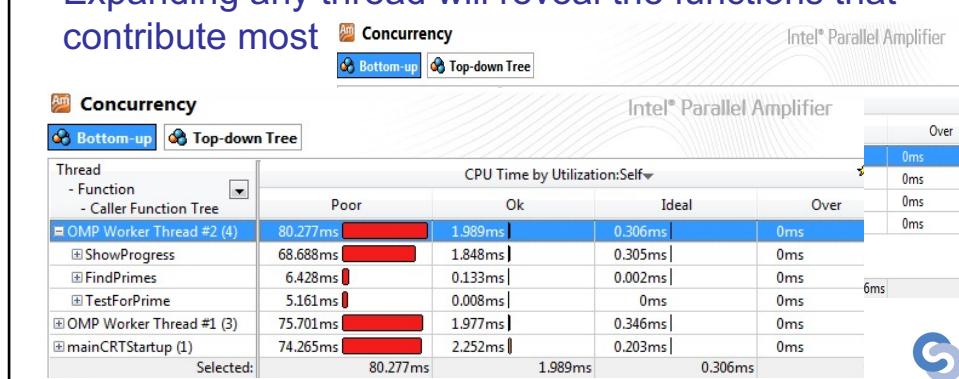
- Function - Caller Function Tree
- ShowProgress is called from FindPrimes and represent the biggest reason concurrency is poor



289

## Parallel Amplifier - Concurrency

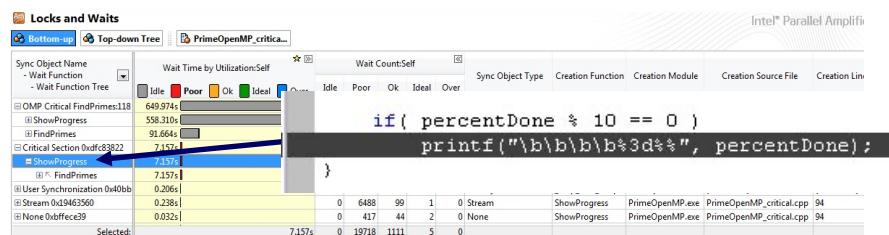
- Thread –Function –Caller Function Tree
- This view shows how each thread contributes to the concurrency issue
- Expanding any thread will reveal the functions that contribute most



290

# Performance

- Double Click ShowProgress in second largest critical section
  - This implementation has implicit synchronization calls - printf
  - This limits scaling performance due to the resulting context switches



## Back to the design stage